Treasure Onwuchekwa Chijioke 23120112046

Answer:

1. Virtual memory paging and swapping mechanisms (Node 3):

Node 3 will be the best choice for this because it has the highest amount of RAM. Reducing the reliance on Virtual memory. This will minimize the reliance on virtual memory paging and swapping.

2. GPU Architecture and Memory Interface (Node 3):

Node 3 will be the best choice for this as it's using an NVIDIA GeForce RTX 3080. This is a high bandwidth GPU with its GDDR6 memory and advanced technology.

3. NPU Influence (Node 3):

Node 3 is the best choice for this as it has NVIDIA Tensor Cores, that are specialized for AI workloads. Good for calculations in machine learning.

4. Power Consumption and Power Management (Node 4):

Node 4 is the best choice for this as it has some really low power specs i.e. Its Celeron J1900 processor and Intel HD Graphics 400 allow for a more power efficient setup.

5. Scalability and Upgrade Path (Node 3):

Node 3 is the best choice for this as it is equipped with the Intel Z390 motherboard that allows for 128 GB of RAM and is compatible with a range of high-performance GPUs and NPUs, making it the most scalable node.

6. Intel/AMD/Celeron Motherboard Constraints (Node 2):

Node 2 is the best choice for this has an AMD 990FX motherboard and AMD FX-8350 processor, that's good for mixed workloads, especially CPU-intensive tasks.